

SRIKANTH REDDY NANDIREDDY

srikanthreddynandireddy@gmail.com | linkedin.com/in/srikanth-reddy-nandireddy
github.com/GodVilan | srikanthreddynandireddy.me | United States (open to Relocation/ Remote)

PROFESSIONAL SUMMARY

AI Engineer building LLM and multi-agent systems: ReAct agents, retrieval-augmented generation (RAG), Model Context Protocol (MCP) servers and clients, and multi-agent orchestration across Azure OpenAI, Anthropic, and AWS. Ships injection-resistant, cost-controlled LLM pipelines with deterministic guardrails, pass^k evaluation harnesses, and calibrated LLM-as-judge scoring, backed by end-to-end MLOps in FastAPI, Docker, MLflow, and GitHub Actions CI/CD. Ranked top 1% (16 of 1,773) in the HackerRank Orchestrate agent hackathon; M.S. Data Science & AI, 4.0 GPA; certified in Anthropic Model Context Protocol and Microsoft Foundry agent development.

CORE SKILLS

Generative AI & LLM Engineering: ReAct Agents, Multi-Agent Orchestration, Retrieval-Augmented Generation (RAG), Prompt Engineering, Prompt Caching, Model Context Protocol (MCP), Microsoft Agent Framework, Function/Tool Calling, Structured Output Enforcement, Critic/Verifier Loops, Multi-Hop Query Decomposition, Vision-Language Models (VLM), Guardrails, Prompt-Injection Defense, Azure OpenAI, Claude API, Google Gemini API

LLM Evaluation & Observability: pass^k Harnesses, LLM-as-Judge (calibrated), Golden Datasets, Ablation Studies, Groundedness & Citation Checks, Confusion Matrices, Macro-F1, Calibration Curves, Token Cost & Latency Tracking, MLflow Experiment Tracking, Regression Test Suites

Retrieval & Vector Search: FAISS (IndexFlatIP), BGE-large-en Embeddings, Sentence Transformers, Hybrid BM25 + TF-IDF Retrieval, Corrective RAG, Chunking & Provenance Tagging

Deep Learning & NLP: PyTorch, TensorFlow, HuggingFace Transformers, BERT, DistilBERT, Wav2Vec2, ViT, Cross-Modal Transformers, Fine-Tuning, Mixed-Precision Training, Tokenization, TF-IDF, Word Embeddings

Machine Learning: Scikit-learn, XGBoost, Random Forest, Logistic Regression, K-Means, Feature Engineering, Class Imbalance Handling, Threshold Optimization, PR-AUC / Precision-Recall Evaluation, Cross-Validation

MLOps & Cloud: FastAPI, Docker, MLflow Model Registry, GitHub Actions, CI/CD, REST APIs, AWS (EC2, S3, SageMaker, Glue, Athena, RDS), Azure OpenAI, Microsoft Foundry, PostgreSQL, Nginx

Data & Analysis: Pandas, NumPy, SciPy, Statsmodels, EDA, Time-Series Forecasting, Hypothesis Testing

Programming & Visualization: Python, SQL (MySQL, PostgreSQL), Matplotlib, Seaborn, Plotly, Dash, Streamlit

AI/ML PROJECTS

Multi-Modal Evidence Review – VLM Claims-Adjudication Agent | HackerRank Orchestrate | [GitHub](#) June 2026

- Ranked 16 of 1,773 (top 1%) in a 24-hour agent hackathon by engineering a two-stage VLM claims-adjudication agent with pixel-to-verdict channel separation, a 14-column structured output contract, and content-hash caching.
- Raised valid-image classification accuracy to 95%, a 35-point gain over a single-pass baseline, by separating perception from reconciliation across two Claude Sonnet calls with deterministic schema and invariant enforcement.
- Blocked 7 of 8 adversarial prompt-injection attacks at roughly \$0.017 per claim by making claim text and EXIF metadata structurally incapable of altering a verdict and prompt-caching static system prompts across calls.

MedsGuard – LLM Medication Agent with a Deterministic Safety Spine | [GitHub](#) 2026

- Eliminated unauthorized-write risk in an LLM medication agent by routing every high-stakes write through an HMAC-signed, single-use, 300-second confirm token over an MCP server on stdio JSON-RPC, never through model output.
- Surfaced 2 reflected prompt-injection leaks that a single-run test missed by executing a pass^k harness with a calibrated LLM-as-judge, then rewrote the parser to emit structured nulls and template-only clarification text.

arXiv ReAct Agent – Research Retrieval Assistant | [GitHub](#) 2026

- Outperformed a BM25 baseline by 7.3x on answer relevance (0.912 vs. 0.125) across 150 papers and 100 QA pairs by building a ReAct agent on BGE-large-en 1024-dim embeddings with FAISS IndexFlatIP vector retrieval.
- Delivered a research assistant exposing 5 agent tools (semantic search, BM25 fallback, live arXiv fetch, summarize, compare) with multi-hop query decomposition and a self-critique refinement loop over persistent workspaces.

TalentFabric AI – 5-Agent Enterprise Readiness Platform | Agents League Hackathon | [GitHub](#) 2026

- Orchestrated a 5-agent Microsoft Agent Framework workflow over 12 learners and 3 teams on Azure OpenAI and Microsoft Foundry, reallocating study hours via a Critic/Verifier loop-back bounded at 2 iterations per learner.
- Fused hybrid BM25 and TF-IDF retrieval with a resilient Microsoft Learn MCP client (24-hour disk cache, 3-tier fallback), logging 6 grounding and accuracy metrics per run to MLflow for release-over-release comparison.

Real-Time Credit Card Fraud Detection – MLOps System | [GitHub](#)

2025

- Deployed a Dockerized FastAPI scoring service on Amazon EC2 serving real-time and batch REST endpoints over 284,807 transactions, automated end to end by a GitHub Actions CI/CD pipeline with API-key auth and rate limiting.
- Versioned 3 fraud classifiers (Logistic Regression, Random Forest, XGBoost) in the MLflow Model Registry with artifacts in Amazon S3, evaluating on PR-AUC because the positive class covers only 0.172% of all rows.

PROFESSIONAL EXPERIENCE

Graduate Laboratory Assistant – Data Science | University of Central Missouri | Warrensburg, MO Feb 2025 – May 2026

- Designed and delivered 15+ machine learning lab sessions per semester to 40+ graduate students, covering Python supervised learning, scikit-learn pipelines, and end-to-end evaluation from data ingestion through prediction.
- Accelerated student proficiency in EDA and model-evaluation workflows using Pandas, NumPy, and Matplotlib, while sustaining a 4.0 GPA across concurrent full-time graduate coursework in deep learning, NLP, and statistics.

Data Analytics Virtual Internship | AICTE EduSkills – Supported by AWS Academy | Remote May 2023 – Jul 2023

- Engineered a serverless cloud analytics pipeline across Amazon S3, AWS Glue, and Athena, automating metadata cataloging and SQL querying over 10K+ semi-structured records without provisioning any persistent compute.
- Built predictive models in Python and Amazon SageMaker to surface dataset trends and feature relationships, completing a structured 10-week AWS Academy data analytics curriculum on cloud-hosted training workflows.

Data Scientist Intern | Swecha | Hyderabad, India May 2022 – Jun 2022

- Improved NLP model reproducibility by 15% across all deployment environments by containerizing Telugu speech-to-text pipelines with Docker, eliminating environment-specific runtime failures at engineering handoff.
- Fine-tuned deep learning speech-to-text models with PyTorch, TensorFlow, and NLTK on dialectally diverse Telugu audio spanning 5+ regional variants, enabling voice-command recognition for a low-resource language.

EDUCATION

M.S. – Data Science & Artificial Intelligence | GPA: 4.0 / 4.0 Aug 2024 – May 2026

University of Central Missouri | Warrensburg, MO
Coursework: Machine Learning, Artificial Intelligence, Neural Networks & Deep Learning, NLP, Statistical Foundations of DS & AI, Advanced Data Visualization

B.S. – Computer Science & Engineering (Honours) | GPA: 9.06 / 10.0 | First Class with Distinction Nov 2020 – Apr 2024

CMR College of Engineering & Technology | Hyderabad, India

CERTIFICATIONS

- Microsoft Applied Skills – Get started developing agents in Microsoft Foundry | Microsoft | Jun 2026
- AI Fluency: Framework & Foundations | Anthropic | Jun 2026
- Introduction to Model Context Protocol | Anthropic | Apr 2026
- AWS Academy Graduate – Data Analytics & Cloud Foundations | Amazon Web Services | Jun 2023
- Foundations: Data, Data, Everywhere (Google Data Analytics Certificate) | Google / Coursera | Feb 2023

HONORS & AWARDS

- Finalist, IICC National Coding Competition (top 1% of 100,000+ participants)
- 5-star HackerRank | 4-star CodeChef in Data Structures & Algorithms
- Graduate Student Achievement Award, University of Central Missouri
- GeeksforGeeks POTD streak, 255 consecutive days of algorithmic problem solving